

# Markov Decision Processes

## Reinforcement Learning

Daniel Parra

24-Dic-2020

## Markov Property

A state  $S_t$ , is Markov if and only if:

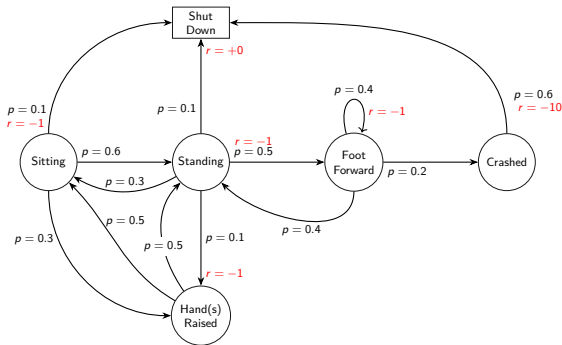
$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, \dots, S_t] \quad (1)$$

### Definition

Formally, for a state  $S_t$  to be **Markov** the probability of the next state  $S_{t+1}$  being  $s'$  should only be dependent on the current state  $S_t = s_t$ , and not on the rest of the past states.



## Markov Reward Process (MRP)





### Markov Reward Process (MRP)

An MRP is defined by  $(s, p(s'|s), r(s), \gamma)$ , where  $s$  are states,  $p(s'|s)$  is the state-transition probability,  $r(s)$  is the reward given  $s$ , and  $\gamma$  is the discount factor.


$$p(s'|s) = \mathbb{P}\{S_t = s' | S_{t-1} = s\} \quad (3)$$

$$r(s) = \mathbb{E}[R_t | S_{t-1} = s] \quad (4)$$

# Markov Decision Process



# Mathematical Formulation



## Dynamics Function

$$p(s', r | s, a) \doteq$$

$$\mathbb{P} \left\{ S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a \right\} \quad (5)$$

$$\sum_{s' \in S} \sum_{r \in R} p(s', r | s, a) = 1 \quad (6)$$

## Some probabilities

### State Transition Probabilities

$$\begin{aligned} p(s'|s, a) &\doteq \mathbb{P} \{ S_t = s' \mid S_{t-1} = s, A_{t-1} = a \} \\ &= \sum_{r \in R} p(s', r | s, a) \end{aligned} \tag{7}$$

## Some probabilities

### Expected rewards for State-Action Pairs

$$\begin{aligned}
 r(s, a) &\doteq \mathbb{E}[R_t | S_{t-1} = s, A_{t-1} = a] \\
 &= \sum_{r \in R} r \cdot \mathbb{P}\{R_t = r | S_{t-1} = s, A_{t-1} = a\} \\
 &= \sum_{r \in R} r \cdot \sum_{s' \in S} \mathbb{P}\{R_t = r, S_t = s' | S_{t-1} = s, A_{t-1} = a\} \\
 &= \sum_{r \in R} r \cdot \sum_{s' \in S} p(s', r | s, a)
 \end{aligned} \tag{8}$$



## Some probabilities

Expected rewards for State-Action-Next-State Triple

$$\begin{aligned}
 r(s, a, s') &\doteq \mathbb{E} \left[ R_t | S_{t-1} = s, A_{t-1} = a, S_t = s' \right] \\
 &= \sum_{r \in R} r \cdot \mathbb{P} \left\{ R_t = r | S_{t-1} = s, A_{t-1} = a, S_t = s' \right\} \\
 &= \sum_{r \in R} r \cdot p(r|s, a, s')
 \end{aligned} \tag{9}$$

**Product Rule applied on Dynamics Function**

$$p(s', r|s, a) = p(s'|s, a) \cdot p(r|s, a, s') \tag{10}$$

$$r(s, a, s') = \sum_{r \in R} r \cdot p(r|s, a, s') = \sum_{r \in R} r \cdot \frac{p(s', r|s, a)}{p(s'|s, a)} \tag{11}$$

# Expected Return

## Formal Definition

$$G_t \doteq R_{t+1} + R_{t+2} + R_{t+3} + \cdots + R_T \quad (12)$$

## Discount

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \quad (13)$$

## Episodic Tasks

$$G_T \doteq \sum_{K=t+1}^T \gamma^{K-t-1} R_K \quad (14)$$

## Policies

### State-Value Function for Policy $\pi$

$$v_{\pi} \doteq \mathbb{E}_{\pi} [G_t | S_t = s] = \mathbb{E}_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s \right], \quad \forall s \in S \quad (15)$$

### Action-Value Function for Policy $\pi$

$$q_{\pi} \doteq \mathbb{E}_{\pi} [G_t | S_t = s, A_t = a] = \mathbb{E}_{\pi} \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \middle| S_t = s, A_t = a \right] \quad (16)$$

## Policies

Give an equation for  $v_\pi$  in terms of  $q_\pi$  and  $\pi$

$$\begin{aligned}
 v_\pi(s) &= \mathbb{E}_\pi [G_t | S_t = s] \\
 &= \sum_{g_t} p(g_t | S_t = s) g_t \\
 &= \sum_{g_t} \sum_a p(g_t, a | S_t = s) g_t \\
 &= \sum_a p(a | S_t = s) \times \sum_{g_t} p(g_t | S_t = s, A_t = a) g_t \quad (17) \\
 &= \sum_a p(a | S_t = s) \times \mathbb{E}_\pi [G_t | S_t = s, A_t = a] \\
 &= \sum_a \underbrace{\pi(a | s)}_{\text{Prob. to take action i.e. policy}} \times \underbrace{q_\pi(s, a)}_{\text{Action-Value Function}}
 \end{aligned}$$

## Policies

Give an equation for  $q_\pi$  in terms of  $v_\pi$  and the four argument p

$$\begin{aligned}
 q_\pi(s, a) &= \mathbb{E}_\pi [G_t | S_t = s, A_t = a] \\
 &= \sum_{s', r, g_t} p(s', r, g_t | s, a) \cdot g_t \\
 &= \sum_{s', r, g_t} p(s', r, g_t | s, a) \cdot [r + \gamma g_{t+1}] \\
 &= \sum_{s', r} p(s', r | s, a) \sum_{g_t} p(g_t | s, a) \cdot [r + \gamma g_{t+1}] \\
 &= \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \sum_{g_t} p(g_t | s') g_t \right] \\
 &= \sum_{s', r} p(s', r | s, a) [r + \gamma \mathbb{E}_\pi [G_t | S_t = s']] \\
 &= \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]
 \end{aligned} \tag{18}$$


## Bellman Equation Derivation

$$\begin{aligned}
 v_{\pi}(s) &\doteq \mathbb{E}_{\pi} [G_t | S_t = s] \\
 &= \mathbb{E}_{\pi} \left[ R_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k R_{(t+1)+k+1} \middle| S_t = s \right] \\
 &= \mathbb{E}_{\pi} [R_{t+1} + \gamma G_{t+1} | S_t = s] \\
 &= \mathbb{E}_{\pi} [R_{t+1} + \gamma \mathbb{E}_{\pi} [G_{t+1} | S_{t+1}] | S_t = s] \\
 &= \mathbb{E}_{\pi} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s] \\
 &= \sum_{s', a, r} p(a, r, s' | s) [r + \gamma v_{\pi}(s')] = \sum_a \pi(a | s) \sum_{r, s'} p(r | s, a) p(s' | s, a) [r + \gamma v_{\pi}(s')] \quad (19) \\
 &= \sum_a \pi(a | s) \sum_r p(r | s, a) r + \gamma \sum_a \pi(a | s) \sum_{s'} p(s' | s, a) v_{\pi}(s') \\
 &= \sum_a \pi(a | s) \sum_{s'} \sum_r p(r, s' | s, a) r + \gamma \sum_a \pi(a | s) \sum_{s'} \sum_r p(s', r | s, a) v_{\pi}(s') \\
 &= \sum_a \pi(a | s) \sum_{s'} \sum_r p(r, s' | s, a) [r + \gamma v_{\pi}(s')]
 \end{aligned}$$

Tomado de: [https://jmichaux.github.io/\\_notebook/2018-10-14-bellman/](https://jmichaux.github.io/_notebook/2018-10-14-bellman/)

## Bellman Equation

**Exercise 3.18** The value of a state depends on the values of the actions possible in that state and on how likely each action is to be taken under the current policy. We can think of this in terms of a small backup diagram rooted at the state and considering each possible action:




Give the equation corresponding to this intuition and diagram for the value at the root node,  $v_\pi(s)$ , in terms of the value at the expected leaf node,  $q_\pi(s, a)$ , given  $S_t = s$ .

$$\begin{aligned} v_\pi(s) &= \sum_a \pi(a|s) q_\pi(s, a) \\ &= \mathbb{E}_\pi [q_\pi(s, a)|s] \end{aligned} \quad (20)$$

Then give a second equation in which the expected value is written out explicitly in terms of  $\pi(a|s)$  such that no expected value notation appears in the equation.

## Bellman Equation

**Exercise 3.19:** The value of an action,  $q_\pi(s, a)$ , depends on the expected next reward and the expected sum of the remaining rewards. Again we can think of this in terms of a small backup diagram, this one rooted at an action (state-action pair) and branching to the possible next states:





## Optimal Policies and Optimal Value Functions

There is always at least one policy that is better than or equal to all other policies, *optimal policy* ( $\pi_*$ ). There may be more than one optimal policy, and they share the same *state-value function* ( $v_*$ ).

$$v_*(s) \doteq \max_{\pi} v_{\pi}(s) \quad \forall s \in S \quad (21)$$

Optimal policies also share the same *optimal action-value function*, denoted  $q_*$ , and defined as:

$$\begin{aligned} q_*(s, a) &\doteq \max_{\pi} q_{\pi}(s, a) \quad \forall s \in S \text{ and } a \in A(s) \\ &= \mathbb{E} [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \end{aligned} \quad (22)$$

## Optimal Value Function

$$\begin{aligned}
 v_*(s) &= \max_{a \in A(s)} q_{\pi_*}(s, a) \\
 &= \max_a \mathbb{E}_{\pi_*} [G_t | S_t = s, A_t = a] \\
 &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\
 &= \max_a \mathbb{E} [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\
 &= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]
 \end{aligned} \tag{23}$$

$$\begin{aligned}
 q_*(s, a) &= \mathbb{E} \left[ R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') | S_t = s, A_t = a \right] \\
 &= \sum_{s', r} p(s', r | s, a) \left[ r + \gamma \max_{a'} q_*(s', a') \right]
 \end{aligned} \tag{24}$$

## Practical Quizz Value Functions and Bellman Equations

**Q1:** A policy is a function which maps **States to probability distributions over actions**.

**Q2:** The term “backup” most closely resembles the term **Update** in meaning.

**Q3:** At least one deterministic optimal policy exists in every Markov decision process. **True** Let's say there is a policy  $\pi_1$  which does well in some states, while policy  $\pi_2$  does well in others. We could combine these policies into a third policy  $\pi_3$ , which always chooses actions according to whichever of policy  $\pi_1$  and  $\pi_2$  has the highest value in the current state.  $\pi_3$  will necessarily have a value greater than or equal to both  $\pi_1$  and  $\pi_2$  in every state! So we will never have a situation where doing well in one state requires sacrificing value in another. Because of this, there always exists some policy which is best in every state. This is of course only an informal argument, but there is in fact a rigorous proof showing that there must always exist at least one optimal deterministic policy.

## Practical Quizz Value Functions and Bellman Equations

**Q4:** The optimal state-value function: **Is unique in every finite Markov decision process.** The Bellman optimality equation is actually a system of equations, one for each state, so if there are  $N$  states, then there are  $N$  equations in  $N$  unknowns. If the dynamics of the environment are known, then in principle one can solve this system of equations for the optimal value function using any one of a variety of methods for solving systems of nonlinear equations. All optimal policies share the same optimal state-value function.

**Q5:** Does adding a constant to all rewards change the set of optimal policies in episodic tasks? **Yes, adding a constant to all rewards changes the set of optimal policies.** Adding a constant to the reward signal can make longer episodes more or less advantageous (depending on whether the constant is positive or negative).

**Q6:** Does adding a constant to all rewards change the set of optimal policies in continuing tasks? **No, as long as the relative differences between rewards remain the same, the set of optimal policies is the same.** Since the task is continuing, the agent will accumulate the same amount of extra reward independent of its behavior.

## Practical Quizz Value Functions and Bellman Equations

**Q7:** Select the equation that correctly relates  $v_*$  to  $q_*$ . Assume  $\pi$  is the uniform random policy.

$$v_* = \max_a q_*(s, a) \quad (25)$$

**Q8:** Select the equation that correctly relates  $q_*$  to  $v_*$ , using four argument function  $p$ .

$$q_*(s, a) = \sum_{s', r} p(s', r|a, s) [r + \gamma v_*(s')] \quad (26)$$

**Q9:** Write a policy  $\pi_*$  in terms of  $q_*$

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } a = \max_{a'} q_*(s, a') \\ 0 & \text{else} \end{cases} \quad (27)$$

The probability of taking an action is constrained between 0 and 1. The value of an action can be arbitrary.

**Q10:** Give an equation for some  $\pi_*$  in terms of  $v_*$  and the four argument  $p$ .

$$\pi_*(a|s) = \begin{cases} 1 & \text{if } v_*(s) = \sum_{r, s'} p(s', r|s, a) [r + \gamma v_*(s')] \\ 0 & \text{else} \end{cases} \quad (28)$$

## Second Quiz

**Q1:** A function which maps \_\_\_ to \_\_\_ is a value function.

- ▶ State-action pairs to expected returns.
- ▶ States to expected returns.

**Q2:** Consider the continuing Markov decision process shown below. The only decision to be made is in the top state, where two actions are available, left and right. The numbers show the rewards that are received deterministically after each action. There are exactly two deterministic policies,  $\pi_{\text{left}}$  and  $\pi_{\text{right}}$

**Q3:** Every finite Markov decision process has:

- ▶ A deterministic optimal policy
- ▶ A unique optimal value function

**Q4:** The \_\_\_ of the reward for each state-action pair, the dynamics function  $p$ , and the policy  $\pi$  is \_\_\_ to characterize the value function  $v_{\pi}$ . **Mean; sufficient**

**Q5:** The Bellman equation for a given a policy  $\pi$ :

- ▶ Expresses state values  $v(s)$  in terms of state values of successor states.

## Second Quiz

**Q6:** An optimal policy: **Is not guaranteed to be unique, even in finite Markov decision processes.**

**Q7:** The Bellman optimality equation for  $v_\pi$ :

- ▶ Holds for the optimal state value function.
- ▶ Expresses state values  $v_*(s)$  in terms of state values of successor states.
- ▶ ~~Holds when  $v_* = v_\pi$  for a give policy  $\pi$~~
- ▶ ~~Expresses the improved policy in terms of the existing policy~~
- ▶ ~~Holds when the policy is greedy with respect to the value function.~~

**Q8:** Give an equation for  $v_\pi$  in terms of  $q_\pi$  and  $\pi$

$$v_\pi(s) = \sum_a \pi(a|s) q_\pi(s, a) \quad (29)$$

**Q9:** Give an equation for  $q_\pi$ , in terms of  $v_\pi$  and the four argument  $p$

$$q_\pi(s, a) = \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')] \quad (30)$$

## Second Quiz

**Q10:** Let  $r(s, a)$  be the expected reward for taking action  $a$  in state  $s$ . Which of the following are valid ways to re-express the Bellman equations, using the expected reward function?

$$\begin{aligned}
 q_{\pi}(s, a) &= r(s, a) + \gamma \sum_{s', a} p(s' | s, a) \pi(a' | s') q_{\pi}(s', a') \\
 q_{*}(s, a) &= r(s, a) + \gamma \sum_{s'} p(s' | s, a) \max_{a'} q_{*}(s', a') \\
 v_{\pi}(s) &= \sum_a \pi(a | s) \left[ r(s, a) + \gamma \sum_{s'} p(s' | s, a) v_{\pi}(s') \right] \\
 v_{*}(s) &= \max_a \left[ r(s, a) + \gamma \sum_{s'} p(s' | s, a) v_{*}(s') \right]
 \end{aligned} \tag{31}$$

**Q11:** Consider an episodic MDP with one state and two actions (left and right). The left action has stochastic reward 1 with probability  $p$  and 3 with probability  $1 - p$ . The right action has stochastic reward 0 with probability  $q$  and 10 with probability  $1 - q$ . What relationship between  $p$  and  $q$  makes the actions equally optimal?

$$7 + 2p = 10q \tag{32}$$